

Deadlocks

M1 MoSIG : Operating Systems

Rouby Thomas

18/11/2014

1 Example

s = 1; available = N; occupied = 0;	
T1	T2
wait(s)	wait(s)
(*) wait(available)	wait(occupied)
.	.
.	.
.	.
post(occupied)	post(available)
post(s)	post(s)

T1 holds s and wait for available. Available will be release by T2 which waits for s. If the thread T1 waits in (*) instruction, there will be a deadlock.

In general, a deadlock is a situation in which each thread in a set of threads hold a resource and is waiting for a resource held by another thread. It is only possible if :

- mutual exclusion : a resource held by a thread cannot be acquired by another.
- no preemption : a resource cannot be taken, even temporarily, from a thread which holds it.
- hold and wait : all the threads in the set hold at least one resource and wait for another one.
- circular wait : t1 waits for t2 to release its resources. t2 waits for t3, t3 waits for t4... tn waits for t1 .

To avoid deadlocks, there are several strategies :

- manage to avoid one of the conditions :
 - share resources so that all threads can access concurrently to them. But it is not always possible.
 - preemption : possible with CPU, but still not always possible.
 - hold and wait : we could require from threads to acquire at once all the resources they need. But it is not practical, not efficient.
 - circular wait : enforce an order for acquiring resources. But it is too difficult in dynamic setups.
- detect deadlocks.
- prevent deadlocks.

2 Resource allocation graph

We need some internal representation of the current state of the system. We are in a situation in which there are mutual exclusion, no preemption and hold and wait => We need to know if circular wait is present/possible.

Model : resources & threads are vertices of a graph.

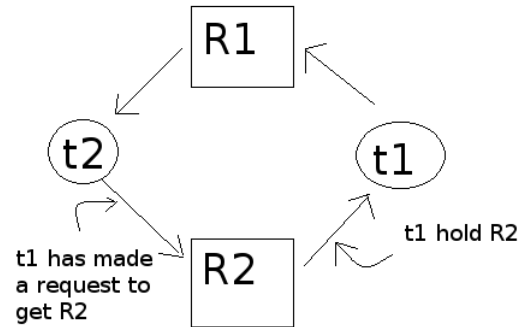


Figure 1: Resource allocation graph

There is a circular wait in the system if there is a cycle in this graph. => Detecting a deadlock is easy using this graph (using, for instance, a bellman-ford algorithm).

Another way to use this graph is to perform a cycle detection for each new request and to decide to refuse it if this induces a cycle in the graph. This model is not suited if resources are divided into types and if threads ask for one instance of a given resource. It can be extended :

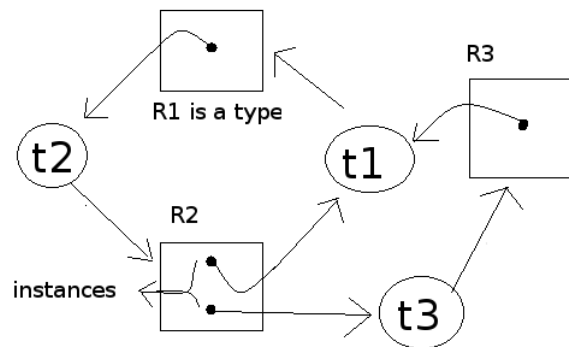


Figure 2: Extended resource allocation graph

In this model, the presence of a cycle doesn't necessarily means that there is a deadlock, only a set of cycles that saturates all the resources it contains, will be a deadlock. -- > Not practical.

3 Allocation matrices and banker algorithm

3.1 Hypothesis

At first we assume that we know :

- a vector Available such that : $\text{Available}[i] = \text{number of instances of resources } i \text{ that are available.}$
- a matrix Allocation such that : $\text{Allocation}[i][j] = \text{number of resources of type } i \text{ allocated to thread } j.$
- a matrix Max such that : $\text{Max}[i][j] = \text{maximal number of resources of type } i \text{ the thread } j \text{ will need.}$
- a matrix Need such that : $\text{Need}[i][j] = \text{Max}[i][j] - \text{Allocation}[i][j]$

We will try to find if there is an order that make the execution of all threads possible :

1. $\text{Work} = \text{Available}$
2. $\text{Finished}[j] = \text{false}$ for all threads.
3. while there is some thread j such that :
 $\text{Finished}[j] = \text{false} \ \&\& \ \text{Need}[i][j] \leq \text{Work}[i] \ \forall,$
then : $\text{Work}[i] += \text{Allocation}[i][j] \ \forall i ;$
 $\text{Finished}[j] = \text{true};$
4. if there is a j such that $\text{Finished}[j] = \text{false}$, the system is not in a safe state.

This previous algorithm finds out if the system is in a safe state or not.

For deadlock prevention, we use the banker algorithm : when an allocation request is issue to the system :

1. Pretend the request is granted :
 $\text{Available} -= \text{request};$
 $\text{Allocation}[*][j] += \text{request};$
2. run the safe state detection
3. if the state is not safe, rollback and deny the request.

3.2 Example of the banker algorithm

	Allocation	A B C	Available	Need	A B C
	0	1 0 1	(0)1 0 2	0	0 2 0
Threads	1	0 0 1		1	2 1 0
	2	0 2 0		2	0 0 1
	3	(2)1 0 0		3	(1)2 2 0
	4	2 0 0		4	2 2 3

Assume that thread 3 requests 1 A resource (1,0,0) (represented in the example with the new numbers between paranthesis)

1. Pretend the request is accepted
2. Run the algorithm to know if the state is safe : $\text{Work} = (0,0,2).$

Thread	Finished
0	f
1	f
2	f
3	f
4	f

Thread 2 has needs \leq Work : Work = (0,2,2).

Thread 0 has needs \leq Work : Work = (1,2,3).

Thread 3 has needs \leq Work : Work = (3,2,3).

Thread 4 has needs \leq Work : Work = (5,2,3).

Thread 1 has needs \leq Work : Work = (5,2,4).

So now we get :

Thread	Finished
0	t
1	t
2	t
3	t
4	t

3. Finished is filled with "true" \Rightarrow the request can be granted.

Assume that thread 4 requests 1A resource and 1C resource (1,0,1).

1. if the request is $>$ Need(thread), or $>$ Available, deny it.
2. Pretend the request is accepted.
3. run the algorithm to know if the state is safe : Work =(0,0,1)

Thread 2 has needs \leq Work : Work = (0,2,1).

Thread 0 has needs \leq Work : Work = (1,2,2).

Thread 3 has needs \leq Work : Work = (4,2,3).

Thread 4 has needs \leq Work : Work = (5,2,3).

Thread 1 has needs \leq Work : Work = (5,2,4).

So now we get again :

Thread	Finished
0	t
1	t
2	t
3	t
4	t

4. the request is granted

Running time of the safe state algorithm is $n^2 * m$ where n is the number of threads and m the number of resource types \Rightarrow Costly, repeated for each allocation request.

4 Detecting Deadlocks vs Preventing them

Let run the system without doing anything and look for a deadlock periodically. Using the resource allocation graph, it's the same algorithm to prevent or detect : find a cycle in the graph. using Allocation matrix : we take an optimistic approach, we assume that thread don't need any more resources => we replace the "Need" Matrix with pending requests, then use the safe state detection to find out if there is a deadlock

If a deadlock is detected :

- preempt resources – > not always possible.
- kill some of the threads :
 - the minimal number of threads that restores a safe state ?
=> risk of killing important threads (ex : Window manager, command interpreter...)
 - killing threads that are not related to the system.
 - let the user decide.
 - choose at random...

In most systems, deadlocks are not handled at all. One can find them in :

- some experimental OS
- some debugging environments